

Aurora / Platinum White Paper

A hosted, human-readable home for the living white paper: what this project is trying to prove, how Platinum and the games fit together, why the conformance program exists, and how the narrative is versioned over time.

Current goal: Turn the current draft into a durable v1 release asset for a broad technical and builder audience: keep the main paper readable, connect it directly to shipped evidence and release discipline, maintain a printable PDF with explicit version/date metadata, and keep related-work and reviewer-minded checks part of the same release surface.

WHITE PAPER
VERSION

vo.4.0-draft

UPDATED

2026-05-17

CURRENT RELEASE

1.4.0.1

LANE

development

WHITE PAPER
STATUS

**living white
paper**

BUILT

May 20, 2026

LATEST NOTE

**Production
follow-through,
white paper,
and
conformance
refresh**

White Paper Project Area

Where this work lives in the repo and how it should be maintained.

Living Draft

The current narrative source of truth is `WHITE\_PAPER.md`. It should keep moving with the project instead of becoming a one-off summary.

Project Area

The `white-paper/` directory is the durable project area for supporting material: the working-area README, citation ledger, and versioned white-paper releases.

Release Memory

Meaningful white-paper revisions should be snapshotted under `white-paper/releases/` with the Markdown, PDF, and PDF metadata preserved together so the narrative can be compared over time just like the software.

Hosted Surface

This generated page is the ongoing hosted, human-readable form. It should stay wired into the same documentation family as the project guide, platform guide, player guide, and release dashboard.

Reader Shape

Who this page is for and how it should balance readability with depth.

Broad Technical Reader

Assume curiosity, technical interest, and some builder intuition, but do not assume deep platform, browser, arcade, or conformance-program expertise.

Readable Main Narrative

The white paper should explain the thesis, architecture, release discipline, evidence program, and AI method clearly without becoming a complete internal manual.

Deeper Surfaces Nearby

Hosted guides, dashboards, release notes, and source docs should provide the extra detail so the main paper can stay focused and persuasive.

Hosted Entry Points

The main ways a reader should enter the white-paper story and related project surfaces.

Current White Paper Draft

Jump directly to the current rendered white paper inside this hosted page.

Citation Ledger

See the maintained list of outside ideas, reference families, and methodological influences.

Current White Paper PDF

Open the current lane PDF version of the white paper with explicit version/date metadata.

PDF Metadata

Inspect the generated metadata for the current lane PDF artifact.

Project Guide

Return to the broader hosted project map for Platinum, Aurora, Galaxy Guardians, ingestion, conformance, and release state.

Lane Project Page

Read the shorter public-facing project summary for the current lane.

Conformance Dashboard

Open the live score, confidence, cost, and investment dashboard when the narrative needs the measured readout.

White Paper Project Source

Open the repo-owned white-paper project area in GitHub.

White Paper Markdown Source

Open the current Markdown source directly.

Review Posture

How the white paper should be reviewed as a release surface instead of treated like an unreviewed note.

HTML Matters

The hosted page should read well on desktop and mobile, keep diagrams and screenshots legible, and avoid accidental raw Markdown leakage or repetitive section structure.

PDF Matters

The printable/exportable PDF should carry explicit version/date information and avoid bad page breaks, oversized dark backgrounds, or weak diagram reproduction.

Reviewer Mentality

Reviewers should tighten repetition, question weak claims, verify references, and treat the white paper as part of the release promise rather than as detached commentary.

Evidence Drift Check

The white-paper review spine should also verify preserved-source integrity and catch stale source-path drift in active evidence docs, not only PDF formatting issues.

White Paper Release History

The first durable release pattern for preserving the narrative over time.

vo.1.0 Snapshot

The first seeded snapshot of the white paper in browsable release form.

vo.1.0 Release Notes

Short note describing what the first white-paper release established and what should happen next.

V1 Release Path

The next best steps for turning the current seed into a true v1 white-paper release.

- Tighten the audience voice so the v1 paper reads clearly as a public-facing promotional and methodological document, not only as an internal orientation memo.
- Recover and cite the earlier Karpathy-style assessment directly so the methodology section can point to a precise prior observation instead of only a conceptual placeholder.
- Add selected screenshots or rendered excerpts from the hosted dashboard, project page, and game surfaces so the paper shows the artifacts it talks about.
- Add a compact internal bibliography that points readers to the canonical source docs: project state, ingestion framework, conformance economics, release policy, and application/platform guides.
- Sharpen the historical arc from 1.0.0 launch to 1.2.0 platform framing to 1.4.0 multi-game and conformance maturity so the story reads as a clear progression.
- Decide which claims should be explicitly public and promotional versus which should remain engineering-facing detail, then trim or regroup sections accordingly.
- Connect the v1 paper to one stable release-quality evidence pack so the narrative and the measurable artifacts can be reviewed side by side.
- Keep the hosted HTML page and lane PDF under the same reviewer checklist so presentation quality, page breaks, and visual clarity improve intentionally over time.

- Refresh the related-work log periodically with focused external searches and short relevance notes instead of letting public references accumulate haphazardly.
- Add a concise section on what this project teaches about generative-AI-assisted software practice beyond games, with examples tied to actual repo artifacts.
- Cut the first intentional white-paper v1 snapshot under `white-paper/releases/` once the citations, visuals, and public voice are all coherent enough to stand alone.

Current White Paper Draft

The living white paper as it exists in the repo today, rendered into the same hosted-document family as the other project guides.

Generated from `WHITE_PAPER.md` during build.

Status: living white paper Current draft: `v0.4.0-draft` Date: 2026-05-17 Audience: broad technical readers, interested builders, collaborators, future reviewers, and public-facing project storytelling

This document is the maintained narrative explanation of what this project is doing, why it is being built this way, and how the approach evolves over time. It is intended to be both a promotional piece and a disciplined reminder to the team: the software, the evidence program, the release process, and the generative-AI workflow all matter together.

See also:

- [white-paper/README.md](#)
- [white-paper/CITATION_LEDGER.md](#)
- [white-paper/REVIEW_CADENCE.md](#)

PLATINUM

Arcade Game Platform



Hosts the cabinet shell, runtime, service adapters, and shared contracts that let multiple fixed-screen shooters live on one coherent platform.

Platinum owns

shell framing, picker, splash, au
leaderboards, feedback, replay, r

Aurora today

first playable pack,
already running on Platinum

What comes next

stabilize the Aurora rerelease baseline
then begin the first playable Galaxy Gu

Documentation graphic · editable SVG · keep aligned with `/PLATINUM_ARCHITECTURE_OVERVIEW.md`

Overview / Preamble

This project is not only a browser game repo.

It is a deliberate attempt to build a professional software program around a harder claim:

- that generative AI can help produce serious, well-released software
- that aggressive iteration does not need to mean careless iteration
- that reference-driven quality can be pursued with explicit evidence instead

of vague taste alone

- that local harnesses, release lanes, review packets, and durable artifacts can

keep AI-assisted work honest

Platinum is the reusable browser-arcade host. **Aurora Galactica** is the first shipped application on that host. **Galaxy Guardians** is the second-game proof that the platform, ingestion program, and conformance discipline can grow beyond a single title.

The larger point is not "we used AI to make a game."

The larger point is that we are building a system in which:

- the product is real
- the releases are real
- the documentation is real
- the tests and harnesses are real
- the evidence is real
- and the model-assisted work is useful because it is forced to leave behind

rerunnable artifacts

Section Overview

- **1. Thesis**: what this project is trying to prove.
- **2. Program snapshot**: where Platinum, Aurora, and Galaxy Guardians stand

right now.

- **3. Five-layer operating model**: platform, games, ingestion, harnesses, and

release economics as one program.

- **4. Ingestion strategy**: how external evidence becomes structured game truth.
- **5. Harnessing and conformance**: how we measure quality instead of asserting

it.

- **6. Release discipline**: how dev, beta, and production remain explicit and

professional.

- **7. Generative AI role**: how model work accelerates the project without

replacing evidence.

- **8. Historical evolution**: how the project moved from launch to platform to

multi-game conformance.

- **9. Citation program:** how outside ideas and source recovery work should be

tracked explicitly.

- **10. Related work and internal references:** how this paper stays short

without losing traceability.

- **11. Living-paper policy:** how this white paper should be maintained and

released over time.

How To Read This Paper

This page is meant to be the readable narrative layer, not the whole archive.

- Read this paper for the story, the method, and the architectural shape of the

project.

- Use the diagrams, screenshots, and charts here as anchors, not as the full

evidence pack.

- If you want implementation detail, detailed metrics, or release-by-release

operational context, jump into the linked hosted documentation rather than making this paper carry everything.

Useful deeper surfaces:

- [project-guide.html](#)
- [platinum-guide.html](#)
- [application-guide.html](#)
- [conformance-dashboard.html](#)
- [release-dashboard.html](#)
- [release-notes.html](#)
- [white-paper.pdf](#)

First Draft

1. Thesis

The core thesis of this project is that generative-AI-assisted software can be built aggressively without becoming hand-wavy, fragile, or unprofessional.

That requires a few non-negotiable rules:

- the software must ship as a real public product, not just as a demo
- improvements must be tied to evidence whenever the question is measurable
- platform boundaries must stay explicit so reuse is deliberate rather than

accidental


- release claims must be backed by committed docs, review artifacts, and

rerunnable checks

- model work should increase leverage, but repeated assessment should

increasingly move into local CPU/browser harnesses

This is why Platinum, Aurora, Galaxy Guardians, ingestion, harnesses, scorecards, review packets, and release notes belong in the same story.



The image above is intentionally simple: it reminds the reader that all of the process, evidence, and release discipline in this paper exist in service of a real playable artifact, not only a methodology exercise.

Further detail:

- [application-guide.html](#)
- [release-notes.html](#)

2. Program Snapshot

As of 2026-05-17, the project can be described in one page:

AREA	CURRENT ROLE	WHY IT MATTERS
Platinum	Shipped browser-arcade host platform	Proves that reusable shell, services, lane model, and release discipline can exist without absorbing game-specific truth.
Aurora Galactica	First shipped playable Platinum application	Serves as the strongest current proof that the platform can host a real public game and improve its conformance over time.
Galaxy Guardians	Preview-first second-game and first-class ingestion/conformance target	Proves that the platform and the evidence program can support a second game without simply cloning Aurora.
Ingestion framework	Source-to-structured-evidence pipeline	Keeps new-game and fidelity work anchored in manifests, clips, event logs, waveforms, contact sheets, and provenance.
Harness and conformance system	Scorecards, correspondence checks, dashboards, and gates	Turns quality claims into measurable, reviewable outputs.
Release and economics program	Lane discipline, review packets, docs refresh, local-vs-cloud resource accounting	Makes the project look and behave like a professional release program rather than an endless prototype.

Current maintained metric read:

SCOPE	CURRENT READ	INTERPRETATION
Aurora Galactica	9.2/10 release-quality conformance	Strong shipped baseline with known high-value gaps, especially audio identity and later depth.
Galaxy Guardians reference preview	7.6/10	Credible second-game proof, but intentionally not presented as production-ready parity.
Galaxy Guardians playtest-weighted preview	6.9/10	Useful reminder that a game can have process maturity before it has public-readiness maturity.
Local-vs-cloud resource split	about 92.8% local wall, about 7.4% declared GPU-equivalent wall	Shows the intended operating doctrine: repeated measurement local, model help strategic.

The evidence program also became more concrete in this pass:

- the recovered March-April Neo-Galaga history now has a canonical

representative preserved-source lane

- the cited Galaga/Galaxian source media now live under repo-owned

preserved-source paths instead of stale old-machine download references

- the white-paper review spine now explicitly checks preserved-source drift as

part of release quality

Aurora Galactica

Current Platinum game pack

Core identity

Galaga-family tribute with capture, rescue, stage flow, challenge cadence, scoring, dual-fighter play, challenge stages, and capture/rescue, enemy and boss behavior, Aurora-owned theme, boss, text, and audio sprites, text, and future Game Designer de

Aurora owns

Current role on Platinum

Aurora is the first fully playable pack and the proof that Platinum can host a real game. The current platform objective is to rerelease Aurora cleanly with Platinum framing while keeping the shipped 1.0.2 gameplay contract intact.

Galaxy Guardians

Future Platinum game pack

Project identity

Galaxian-inspired sibling game for Platinum
Separate planning identity from Galaxy Guardians
reference lineage and from Aurora's me

First intended slice

formation rack, independent dives,
flagship/escort scoring, single-shot press
wrap-around threat, no Aurora capture/resc

Current role on Platinum

Right now Galaxy Guardians is a shell preview and coming-soon splash, not a playable
The right next step is to start gameplay only after Aurora's Platinum rerelease ba
is trusted and the platform/game seam remains clear.

These pack views help a broad reader understand one of the project's central claims: Aurora Galactica and Galaxy Guardians are not supposed to be two skins on one game. They are meant to be separate applications living on one host platform.

TODO illustration: Choose a small three-panel progression strip that shows how the public face of the project evolved from 1.0.0 launch to 1.2.0 Platinum framing to 1.4.0 multi-game posture. The most illustrative version may be gameplay first, shell first, or docs/release-surface first, and we should pick that deliberately rather than guessing.

Further detail:

- [project-guide.html](#)
- [platinum-guide.html](#)

3. Five-Layer Operating Model

The repo already describes the work as a layered system. The white paper should make that model legible at a glance.

Reference games, videos,
manuals, and source
artifacts

Ingestion framework

Game-owned evidence:
manifests, windows, event
logs, cue catalogs, stage
summaries

Platinum platform

Applications: Aurora
Galactica and Galaxy
Guardians

Harnesses, scorecards,
correspondence reports,
dashboards

Release lanes, docs,
review packets, and public
claims

Observed gaps, player
feedback, and next
investment choices



Platinum Architectural Overview

Platform layers, game-pack ownership, and the current hosting model for Aurora Galactica.

Platinum Shell

Cabinet frame, marquee, side rails, popup framing, build stamp, picker surfaces, platform splash.

Platinum Runtime and Services

Pack install path, wait mode, pause, scaling, auth/session, leaderboard transport, replay plumbing,

Game Pack Layer

Each game owns its identity, rules, mechanics, assets, audio, content, and game-specific harnesses.

Aurora Galactica: live playable Galaxy Guardians: preview sibling Shared contracts and harnesses

What belongs to Platinum

Shell, runtime, platform splash, picker, service adapters, publish/release checks, platform-only harnesses.

What belongs to a game pack

Stages, enemies, boss behavior, scoring, challenge log, sprites, text, audio identity, designer-owned definitions.

Platinum vs Game Packs

A clearer ownership split between the host platform and the games it carries.

Platinum

Host platform

- shell and framing
- pack picker
- platform splash
- auth and leaderboard adapters
- replay plumbing
- platform-only harnesses
- publish and lane checks

Aurora Galactica

Current playable game pack

- formations and stage flow
- capture and rescue
- challenge stages
- Aurora-owned art, text, audio

Galaxy Guardians

Future sibling game pack

- Galaxian-inspired rules
- its own stage and enemy identity
- its own assets and designer surf

Shared seam

This is where Platinum and packs meet

- metadata and game key
- shell themes
- stage presentation
- challenge layout
- scoring tables
- entity contract and capability flags
- pack-specific harness families

Rule of thumb

If it should work for multiple games, it probably belongs to Platinum.

The important discipline is separation of ownership:

- Platinum owns shell, hosting, shared services, contracts, and release

framing.

- Games own rules, scoring, progression, audiovisual identity, and their own

conformance truth.

- Ingestion owns provenance and structured evidence.
- Harnesses own repeatable evaluation.
- Release discipline owns the public promise.

When those layers blur, the project becomes harder to explain, harder to test, and easier to accidentally fake.

Further detail:

- [project-guide.html#platform-vs-applications](#)
- [platinum-guide.html](#)
- [PROJECT_STATE_AND_CONFORMANCE_PROGRAM.md](#)

4. Ingestion Strategy

Ingestion is the front half of engineering, not a side notebook.

The project does not want new games or fidelity improvements to come mainly from memory, vibes, or post-hoc rationalization. Instead, it wants evidence to arrive in structured forms that can be reused:


- source manifests
- preserved clips and windows
- frame contact sheets
- waveforms and spectrograms
- reference-side event logs
- semantic annotations
- confidence notes
- scorer and correspondence targets

For **Aurora**, this keeps Galaga-like timing, audio, pressure, and stage-shape questions grounded in real artifacts.

For **Galaxy Guardians**, ingestion matters even more. It is the mechanism that prevents the second game from turning into "Aurora with different labels." The game should become more complete by promoting **Galaxian** evidence into game-owned scoring, wave timing, sprite identity, audio expectations, and runtime correspondence checks.

In short:

- external artifacts teach the game
- ingestion turns artifacts into structured evidence
- the application implements against that evidence
- Platinum stays the host rather than the hidden author of the game





These reference contact sheets are useful because they show the project’s ingestion claim in a form a non-expert can understand quickly. We are not only describing classic arcade behavior; we are collecting windows, studying them, and turning them into reusable evidence.

That claim is now easier to defend concretely because the repo carries preserved-source lanes as well as derived analyses. The recovered Galaga source videos, Freesound pack, Galaxian cabinet way, and the first representative historical Neo-Galaga archive slice now live under `reference-artifacts/preserved-sources/` with manifests and checksums. In other words, the project is moving source recovery out of memory and into committed provenance.

TODO illustration: Pick the single best “ingestion in action” image for v1. The strongest option might be a contact sheet, a waveform-plus-contact-sheet pair, or a staged comparison between raw source footage and the structured artifact family that comes out of it.

Further detail:

- project-guide.html#classic-arcade-ingestion-framework-doc
- application-guide.html
- [CLASSIC_ARCADE_INGESTION_FRAMEWORK.md](#)
- reference-artifacts/preserved-sources/README.md

5. Harnessing And Conformance

This project is serious about the difference between "better" and "better by a rerunnable measure."

The conformance system exists so that quality can be described with more precision than a mood:

- timing correspondence
- sequence correspondence
- outcome correspondence
- spatial correspondence
- visual correspondence
- audio correspondence
- persona and progression correspondence

The current Aurora scorecard turns this into a twelve-category quality model. That matters for two reasons.

First, it helps choose investments that are actually player-visible.

Second, it protects the team from false confidence. A 10/10 is explicitly not "perfect"; it means "maxed at current scorer resolution." Better evidence or a better evaluator can lower a score while making the project more truthful.


The harness program also stays intentionally classified:

- `platform` harnesses protect shell, hosting, docs, and shared services
- `application` harnesses protect game-specific rules and behavior
- `boundary` harnesses protect the seam between Platinum and the games

Representative committed commands in this strategy include:

- `npm run harness:measure`
- `npm run review:code`
- `npm run review:ledger`
- `npm run harness:check:galaxy-guardians-first-class-conformance`

This is the deeper quality claim of the project: bugs, polish, and release readiness should increasingly move from memory and opinion into explicit checks, artifacts, and dashboards.



Persona Performance Distribution

Seeded full-run persona gameplay; each line is one generic test persona across 30 current Aurora runs.



The value of these charts is not only that they look rigorous. They show that the project tries to externalize quality questions into surfaces that can be inspected, debated, and rerun.

Further detail:

- conformance-dashboard.html
- project-guide.html#release-conformance-dashboard-doc
- project-guide.html#conformance-economics-doc

6. Release Discipline And Professionalism

The project treats release engineering as part of product quality.

That means the release lanes are not cosmetic:

1. local `localhost`
2. hosted `/dev`
3. hosted `/beta`
4. hosted `/production`

Each lane carries a different stability promise, documentation expectation, and testing posture. The project is intentionally trying to behave like a software program with real public accountability:

- release notes are first-class
- docs are part of the release surface
- review packets and review-learning ledgers are durable evidence
- the white paper must read well both as hosted HTML and as printable PDF

- production claims should come from an approved beta lineage
- major releases should refresh dashboards, scorecards, and strategic docs

This matters because AI-assisted speed is only impressive if the public result still feels trustworthy.

The "reviewer" mentality should therefore be explicit. The paper is not done just because the words are present. The release surface should also be reviewed for:

- repeated ideas that can be tightened
- diagrams or images that create awkward whitespace or weak page breaks
- TODOs that are unclear or more revealing of indecision than of real planning
- HTML reading flow on desktop and mobile
- PDF export quality, especially around image scale, table breaks, and diagram

legibility

Further detail:

- release-dashboard.html
- release-notes.html
- project-guide.html#testing-doc
- project-guide.html#code-review-model-doc
- white-paper/REVIEWER_CHECKLIST.md
- white-paper/REVIEW_CADENCE.md

7. How Generative AI Fits

The project does use generative AI heavily, but not as a substitute for engineering structure.

The intended operating doctrine is:

- use models to design evaluators, synthesize options, write code, review code,

summarize evidence, and tighten the next decision

- use local CPU/browser harnesses for repeated measurement, sweeps, and

regression checks

- convert model-assisted insight into committed repo artifacts whenever

possible

- log resource use and compare spend against score movement
- keep human review, public release notes, and versioned documentation in the

loop


The repo already describes one part of this explicitly as a "Karpathy-loop-like" pattern:

- inspect concrete examples
- improve the dataset and evaluator
- make a small candidate change


- run it
- study failures
- fold the learning back into the system

That is a strong fit for the broader project identity. The point is not merely to ask a model for code. The point is to build a system in which model help leaves behind better evaluators, better artifacts, and cheaper future decisions.

Tracked Compute Minutes By Resource Class



Tracked Wall Minutes Per +1 Score Point



These charts help keep the AI story grounded. The point is not only that model assistance exists; it is that the project is trying to compare that assistance with local repeatable measurement and with visible quality movement.

Further detail:

- project-guide.html#conformance-economics-doc
- [CONFORMANCE ECONOMICS.md](#)

8. Working Loop

The operating loop of this project is more important than any single feature.



This loop explains how the project tries to be both aggressive and controlled. The aggressiveness comes from fast iteration and model-assisted leverage. The control comes from evidence, harnesses, explicit ownership boundaries, and release discipline.

TODO illustration: Add one compact “question -> evidence -> harness -> change -> rerun” visual from a real case study. Audio cue alignment, stage-opening timing, or a Galaxy Guardians reference-promotion slice are the strongest current candidates, but we should choose the one that is most legible to a broad reader.

9. Historical Evolution So Far

The release notes already show a clear arc, and the white paper should make it easy to retell.

RELEASE	MEANING	STRATEGIC SHIFT
1.0.0	First public Aurora launch	The project became a real public product with live scoring, pilot identity, replay visibility, and a real release ladder.
1.2.0	Platinum Release 1	Aurora was reframed as the first application on a reusable platform, making platform/application separation explicit.
1.4.0	Current multi-game and conformance baseline	The public line now carries stronger documentation, review evidence, persona/replay follow-through, and a clearer Galaxy Guardians posture.

This means the project has already moved through three meaningful phases:

1. prove a public game can ship
2. prove the host platform is real
3. prove multi-game and conformance maturity can become part of the release

identity

The next phase should be to prove that this method scales:

- deeper Aurora conformance
- stronger Galaxy Guardians first-class completeness
- cleaner shared Platinum operations
- more reusable ingestion and assessment infrastructure

TODO illustration: Build a release-history gallery with one screenshot or architectural surface per milestone. The current paper names the milestones clearly, but a short visual strip would make the progression easier to absorb at a glance.

Further detail:

- [release-notes.html](#)
- [project-guide.html#release-note-140-beta-1-doc](#)
- [project-guide.html#release-note-130-production-refresh-doc](#)

10. Citation Program

This white paper should not quietly absorb ideas or source recovery work without naming them.

We want a maintained citation program that records:

- what outside work or source material influenced us

- how we used it
- what we learned from it
- what changed in the repo because of it
- what still needs to be recovered, linked, or tightened

The living ledger for that work starts here:

- [white-paper/CITATION_LEDGER.md](#)

The source-recovery side of that program now has a matching repo-owned surface:

- [reference-artifacts/preserved-sources/README.md](#)

That matters because provenance is not only a footnote here. It is a release quality concern. If a timing study, audio comparison, or historical claim depends on a file that only exists in somebody's old downloads folder, the project is less professional than it looks.

The first open citation debt is the prior standalone assessment of the Karpathy-style research/evaluator loop. The repo contains the conceptual thread already, but the older assessment should be recovered and linked directly in a future white-paper release rather than reconstructed from memory.

Further detail:

- [white-paper/CITATION_LEDGER.md](#)

11. Related Work

This project should periodically stop and look outward.

The right pattern is not to stuff the paper with literature. The right pattern is to do focused searches, add high-signal sources, explain their relevance in plain language, and keep the public references linked to a maintained log.

Current seeded related-work set:

- [Anthropic, "Building effective agents" \(2024-12-19\)](#): relevant because it argues for simple, composable agent patterns and evaluator-optimizer loops rather than ornamental workflow complexity.
- [Anthropic, "Writing effective tools for agents - with agents" \(2025-09-11\)](#): relevant because our harnesses, scripts, and release tools are all explicit contracts between model assistance and deterministic system behavior.
- [Anthropic, "Demystifying evals for AI agents" \(2026-01-09\)](#): relevant because it reinforces our investment in transcripts, graders, repeated trials, and measurable evaluator quality instead of anecdotal confidence.
- [Anthropic, "Trustworthy agents in practice" \(2026-04-09\)](#): relevant because it frames guardrails, reviews, and operator-visible controls as part of the product surface, not only as implementation details.
- [METR, "Measuring AI Ability to Complete Long Tasks" \(2025-03-19\)](#): relevant because it gives a sharper lens on what kinds of work current agents can sustain autonomously, which is directly relevant to our preference for small evaluator-visible loops and local reruns over vague claims of full autonomy.
- [OpenAI, "PaperBench" \(2025-04-02\)](#): relevant because it shows how large agentic tasks become more reviewable when decomposed into explicit rubrics and many individually gradable subtasks, which is close to how our own harnesses and conformance categories create reviewable surfaces.

Maintained deeper log:

- [white-paper/RELATED_WORK.md](#)

12. Internal Canonical Docs

This paper should stay readable because the repo already has deeper canonical surfaces nearby.

The shortest list of internal references that best supports the claims here is:

- [project-guide.html](#)
- [platinum-guide.html](#)
- [application-guide.html](#)
- [conformance-dashboard.html](#)
- [release-dashboard.html](#)
- [CLASSIC_ARCADE_INGESTION_FRAMEWORK.md](#)
- [PROJECT_STATE_AND_CONFORMANCE_PROGRAM.md](#)
- [CONFORMANCE_ECONOMICS.md](#)
- [reference-artifacts/preserved-sources/README.md](#)

If the main paper starts to feel long, that is usually a sign that one of these surfaces should carry more of the detail instead.

13. Why This Project Matters

The project matters because it is trying to demonstrate a concrete alternative to two weak extremes.

It is not:

- slow, ceremonial process that kills iteration
- or fast AI-assisted output that cannot explain or defend itself

Instead, it aims for a middle path:

- ambitious shipping
- real public releases
- strong platform boundaries
- evidence-backed improvement
- local-first repeated measurement
- model-assisted leverage
- versioned documentation and release storytelling

If that works, the result is more than a good arcade project. It becomes a useful pattern for how generative AI can participate in professional software work without dissolving quality standards.

This is also why the paper should remain readable. A broad technical reader does not need every source artifact inline. They need a coherent narrative, selected visual proof, and obvious places to go next if they want more depth.

14. Living White Paper Policy

This document should evolve the same way the project evolves: intentionally, versioned, and with historical memory preserved.

Working policy:

- `WHITE_PAPER.md` is the current living draft
- the hosted HTML and printable PDF should be kept aligned as two views of the

same maintained release surface

- meaningful revisions should be snapshotted under `white-paper/releases/`
- each snapshot should preserve the exact white paper text for that release
- when a release PDF exists, the snapshot should preserve the Markdown, PDF,

and generated PDF metadata together

- the citation ledger should be updated when outside work materially changes the

project story

- the related-work log should be refreshed periodically with focused online

searches and brief relevance commentary

- the recurring reviewer rhythm in `white-paper/REVIEW_CADENCE.md` should be

treated as normal maintenance, not as a one-off cleanup exercise

- reviewer-checklist expectations should be treated as part of release quality,

not optional cleanup

- every meaningful software release does not need a new white paper release, but

every strategic narrative shift probably does

Good triggers for a new white paper release:

- a major public release family
- a major shift in the conformance program
- a stronger second-game milestone
- a meaningful change in the AI/harness/evidence operating model
- recovery of an important citation or conceptual influence

Immediate Next Additions

- Recover and link the earlier Karpathy-style assessment if it exists outside

this repo.

- Add one deliberate progression gallery for milestone history and one deliberate

“evidence in action” case-study image once we decide which examples explain the project most clearly.

- Decide whether the historical Neo-Galaga archive should stop at the current

representative 22-pair preserved slice or grow into a larger non-git or secondary preserved-source lane.

- Keep the HTML and PDF release surfaces under reviewer scrutiny so that spacing,

diagrams, repeated ideas, and print behavior all improve with the narrative.

- Keep the audience tuned for a broad technical and builder readership: assume

interest, assume intelligence, but do not assume deep prior expertise.

Citation Ledger

Maintained list of outside ideas, source families, how they were used, and what still needs to be recovered or tightened.

Generated from `white-paper/CITATION_LEDGER.md` during build.

This ledger tracks source families, methodological influences, and outside ideas that materially shape the project story.

The goal is not only to cite things. The goal is to record:

- what influenced the work
- how that influence showed up in the repo
- what the team learned from it
- what citation or linkage debt still remains

Status Key

- `linked`: the reference is represented clearly enough in the repo today
- `partial`: the idea is present, but the exact earlier note, external source,

or final public citation still needs to be recovered or tightened

- `queued`: important enough to track now, but not yet integrated well enough

to claim as a polished citation

Current Entries

REFERENCE	KIND	HOW WE USE IT	WHAT WE LEARNED	CURRENT REPO ANCHOR	STATUS
Karpathy-style evaluator loop and earlier project assessment	conceptual / methodological	Shapes the idea that we should inspect concrete examples, improve evaluators, make small candidate changes, rerun, and study failures instead of tuning only by opinion.	Better evaluators can be as important as better runtime code. A stricter scorer can lower a score while making the project more truthful.	PROJECT_STATE_AND_CONFORMANCE_PROGRAM.md , CONFORMANCE_ECONOMICS.md , RELEASE_NOTE_1.3.0.1_HOSTED_DEV_REVIEW.md	partial
Anthropic, "Building effective agents" (2024-12-19)	external methodological reference	Reinforces the idea that agentic systems should prefer simple, composable loops and explicit evaluator structures instead of ornamental complexity.	Simpler loops become more legible when the evaluator, harness, and release artifacts are visible to reviewers.	WHITE_PAPER.md , white-paper/RELATED_WORK.md , PROJECT_STATE_AND_CONFORMANCE_PROGRAM.md	linked
Anthropic, "Writing effective tools for agents - with agents" (2025-09-11)	external methodological reference	Supports our view that tools are explicit contracts and that agent quality depends heavily on the quality, shape, and reviewability of those tools.	Better tools and better tool descriptions are a form of product quality, not only implementation detail.	WHITE_PAPER.md , white-paper/RELATED_WORK.md , TESTING_AND_RELEASE_GATES.md	linked
Anthropic, "Demystifying evals for AI agents" (2026-01-09)	external evaluation reference	Supports our investment in repeated trials, transcripts, graders, and explicit evaluation design for agent-assisted work.	Eval become more useful when they are cheap to rerun, narrow in scope, and part of the everyday engineering loop.	WHITE_PAPER.md , white-paper/RELATED_WORK.md , CONFORMANCE_ECONOMICS.md	linked
Anthropic, "Trustworthy agents in practice" (2026-04-09)	external governance / operations reference	Aligns with our insistence that guardrails, human review, release notes, and reviewer-visible controls are part of the	Trustworthiness is easier to discuss honestly when it is backed by durable checks and visible	WHITE_PAPER.md , white-paper/RELATED_WORK.md , RELEASE_POLICY.md , CODE_REVIEW_MODEL.md	linked

REFERENCE	KIND	HOW WE USE IT	WHAT WE LEARNED	CURRENT REPO ANCHOR	STATUS
		product and release surface.	operational policy.		
METR, "Measuring AI Ability to Complete Long Tasks" (2025-03-19)	external capability/evaluation reference	Helps explain why the project prefers narrow, rerunnable loops and modest autonomy claims rather than treating all agentic work as equally reliable.	Measuring agent capability by task duration is a useful complement to benchmark-style scores and fits our local-rerun doctrine.	<code>WHITE_PAPER.md</code> , <code>white-paper/RELATED_WORK.md</code> , <code>CONFORMANCE_ECONOMICS.md</code>	linked
OpenAI, "PaperBench" (2025-04-02)	external evaluation/reference-design influence	Supports our instinct to decompose complex AI-assisted work into explicit rubrics, gradable subtasks, and reviewer-visible evidence.	Hard agentic work becomes easier to discuss honestly when the grading structure is explicit instead of implied.	<code>WHITE_PAPER.md</code> , <code>white-paper/RELATED_WORK.md</code> , <code>CODE_REVIEW_MODEL.md</code>	linked
Galaga gameplay footage, manuals, clips, and extracted artifacts	reference corpus	Grounds Aurora timing, audio, stage cadence, visual comparison, and correspondence work in preserved material.	Manual impressions are useful, but clipped windows, event logs, and aligned audio/visual artifacts make fidelity work reviewable and reusable.	<code>reference-artifacts/</code> , <code>VIDEO_ALIGNMENT_PROGRAM.md</code> , <code>CORRESPONDENCE_FRAMEWORK.md</code>	linked
Galaxian gameplay footage and sibling-game source package	reference corpus	Grounds Galaxy Guardians in its own source lineage so the game can become a true sibling application rather than a relabeled Aurora variant.	Ingestion is most valuable when it arrives before design hardens. Second-game credibility depends on game-owned evidence, not borrowed first-game behavior.	<code>CLASSIC_ARCADE_INGESTION_FRAMEWORK.md</code> , <code>APPLICATIONS_ON_PLATINUM.md</code> , <code>CONFORMANCE_METRICS_OVERVIEW.md</code>	linked

REFERENCE	KIND	HOW WE USE IT	WHAT WE LEARNED	CURRENT REPO ANCHOR	STATUS
Recovered old-machine source media and representative Neo-Galaga archive	provenance / evidence discipline	Makes source recovery, historical runs, and cited reference media part of the repo-owned evidence program rather than depending on remembered old download paths.	Provenance is stronger when recovered sources have preserved-source lanes, manifests, hashes, and active-doc links instead of only intake notes.	<code>reference-artifacts/preserved-sources/</code> , <code>reference-artifacts/ingestion/downloads-old-all-2026-05-17/</code> , <code>WHITE_PAPER.md</code>	linked
Review packet and review-learning ledger model	operational discipline	Makes AI-assisted and fast-moving changes reviewable through durable packets, issue categories, and production dispositions.	Review value compounds when repeated findings become harnesses, release checks, or documented non-goals instead of disappearing into chat.	<code>CODE_REVIEW_MODEL.md</code> , <code>REVIEW_LEARNING_LEDGER.md</code>	linked
Local-first compute doctrine for conformance work	operating doctrine	Pushes repeated measurement into local CPU/browser harnesses while reserving model work for strategy, synthesis, evaluator design, and selected analysis.	The project becomes cheaper and more trustworthy when model assistance leaves behind committed local logic and measurable artifacts.	<code>CONFORMANCE_ECONOMICS.md</code> , <code>PROJECT_STATE_AND_CONFORMANCE_PROGRAM.md</code>	linked

Open Citation Debt

- Recover the earlier standalone Karpathy-related assessment and replace the current conceptual placeholder with a precise internal or external citation.
- Decide which outside frameworks deserve public-facing mention in the white paper versus remaining internal working influences.
- Keep a dated related-work refresh cadence so public references are added on purpose rather than only when remembered opportunistically.

Illustration Plan

Working list of selected visuals, candidate deeper assets, and the open illustration choices that still need deliberate discussion.

Generated from `white-paper/ILLUSTRATION_PLAN.md` during build.

This document tracks the visuals that support the white paper, the deeper reference materials that should stay nearby in hosted docs, and the places where we still need to choose the most illustrative image or chart deliberately.

Selection Rules

- Prefer visuals that a broad technical or builder reader can interpret quickly.
- Use the main paper for orientation-quality visuals, not for every supporting artifact.
- Put deeper evidence families in the hosted guides, dashboards, and linked source documents instead of overloading the main narrative.
- When a visual choice is ambiguous, keep a placeholder in the paper and record the decision debt here.

Current In-Paper Visuals

WHITE-PAPER SECTION	CURRENT ASSET	WHY IT WORKS
Overview / thesis	reference-artifacts/diagrams/platinum/platinum-hero.svg	Gives the paper an immediate project identity and platform-level frame.
Thesis	export.mov.png	Shows that the evidence and release program serve a real playable artifact.
Program snapshot	reference-artifacts/diagrams/platinum/aurora-pack-card.svg	Makes Aurora legible as an application on the platform.
Program snapshot	reference-artifacts/diagrams/platinum/galaxy-guardians-pack-card.svg	Shows second-game identity without requiring a long explanation.
Five-layer operating model	reference-artifacts/diagrams/platinum/platinum-platform-stack.svg	Reinforces platform-versus-application separation.
Five-layer operating model	reference-artifacts/diagrams/platinum/platinum-pack-separation.svg	Helps explain ownership boundaries at a glance.
Ingestion strategy	reference-artifacts/analyses/galaga-stage-opening-timing/2026-04-12-main-a777fba/opening-contact-tight.png	Makes ingestion concrete through a visible reference-study artifact.
Ingestion strategy	reference-artifacts/analyses/galaxian-reference/matt-hawkins-arcade-intro/frames/contact-sheet-reference-window.jpg	Supports the claim that Galaxy Guardians is grounded in its own source family.
Harnessing and conformance	reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/score-trends.svg	Turns progress into an at-a-glance measurable story.
Harnessing and conformance	reference-artifacts/analyses/persona-performance-distribution/performance-lines.svg	Shows that quality is evaluated across viewpoints, not through one metric alone.
Release and economics	reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/compute-minutes-by-resource.svg	Makes local-first measurement strategy visible.
Release and economics	reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/cost-per-positive-score-point.svg	Connects release ambition to investment discipline.

Deeper Supporting Visuals

These are better linked from hosted guides or follow-on detail pages than pushed directly into the main narrative unless a specific section needs them.

- [reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/largest-score-deltas.svg](#)
- [reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/gpu-equivalent-use-by-purpose.svg](#)
- [reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/cpu-use-by-purpose.svg](#)
- [reference-artifacts/analyses/conformance-economics/2026-05-14-1c788342/gameplay-improvement-by-project-part.svg](#)
- hosted [conformance-dashboard.html](#)
- hosted [release-dashboard.html](#)
- hosted [project-guide.html](#)

- hosted `public-project-page.html`

Open Illustration TODOs

Release-history progression strip

- White-paper placeholder exists.
- Best candidate still needs discussion:
- `1.0.0` public game surface
- `1.2.0` Platinum framing surface
- `1.4.0` multi-game and conformance surface
- Decision question: should the strip emphasize gameplay evolution, platform

architecture, or public release/documentation maturity?

Ingestion in action

- White-paper placeholder exists.
- Best candidate still needs discussion:
- contact sheet only
- waveform plus contact sheet pair
- raw source clip to structured artifact pipeline graphic
- Decision question: what most clearly teaches the ingestion idea to a

non-expert reader in one glance?

Evidence loop case study

- White-paper placeholder exists.
- Best candidate still needs discussion:
- question -> artifact -> harness -> result graphic
- screenshot pair showing before and after measured correction
- release-note excerpt plus chart plus artifact collage
- Decision question: which single case study best shows how AI-assisted work is

kept honest by rerunnable evidence?

Milestone gallery

- White-paper placeholder exists.
- Best candidate still needs discussion:
- one screenshot per milestone release
- one diagram per maturity phase
- one mixed gallery of game, platform, and dashboard surfaces
- Decision question: should the gallery feel product-led, architecture-led, or

release-led?

Likely Next Additions

- Add one screenshot from the hosted conformance dashboard.
- Add one screenshot from the hosted release dashboard.
- Add one architectural illustration that shows ingestion, games, Platinum, and

release lanes on a single page.

- Add one progression visual based on preserved historical release surfaces.

Related Work Log

Dated log of external references, why they matter here, and when they should be refreshed.

Generated from `white-paper/RELATED_WORK.md` during build.

This log tracks external work worth preserving alongside the white paper.

It is deliberately narrower than a literature survey. The goal is to keep a small, high-signal set of references that help explain why this project looks the way it does.

Working Rule

- Prefer focused, occasional searches over bulk accumulation.
- Prefer primary or close-to-primary sources when the topic is methodological.
- Record why the work matters here, not only that it exists.
- Update the main white paper only with the most relevant current items; keep

the longer memory here.

Current Seed Set

DATE ADDED	WORK	RELEVANCE HERE	LINK
2026-05-16	Anthropic, *Building effective agents* (2024-12-19)	Supports our preference for simple loops, visible evaluator structure, and avoiding ornamental agent complexity.	Source
2026-05-16	Anthropic, *Writing effective tools for agents - with agents* (2025-09-11)	Closely matches our view that harnesses, scripts, and release tools are explicit contracts between model assistance and deterministic behavior.	Source
2026-05-16	Anthropic, *Demystifying evals for AI agents* (2026-01-09)	Reinforces our investment in evals, reruns, transcripts, graders, and narrow measurable questions.	Source
2026-05-16	Anthropic, *Trustworthy agents in practice* (2026-04-09)	Supports the idea that controls, guardrails, human review, and operator-visible release process are product concerns, not only policy concerns.	Source
2026-05-17	METR, *Measuring AI Ability to Complete Long Tasks* (2025-03-19)	Useful because it frames agent capability in terms of task duration and reliability, which helps explain why this project prefers small, reviewer-visible loops and local reruns over stronger autonomy claims.	Source
2026-05-17	OpenAI, *PaperBench* (2025-04-02)	Useful because it shows how hard agentic work becomes more reviewable when decomposed into explicit rubrics and many gradable subtasks, which is close to how our own conformance categories make progress inspectable.	Source

Refresh Triggers

- major white-paper revision
- major release-policy or conformance-process revision
- meaningful new methodological influence worth naming publicly
- significant change in the broader agent/evals/tooling conversation that bears

directly on this project

Reviewer Checklist

Release-minded checklist for reading the white paper critically across narrative, HTML, PDF, and related-work quality.

Generated from `white-paper/REVIEWER_CHECKLIST.md` during build.

This checklist is meant to make the reviewer mentality explicit.

The white paper is part of the release surface. It should be reviewed with the same seriousness as other user-visible documentation and release artifacts.

Core Stance

- Read as a reviewer, not as the author.
- Prefer tighter and clearer over longer and more complete.
- Challenge repetition, overclaiming, weak evidence links, and vague diagrams.
- Treat HTML presentation, PDF presentation, and narrative quality as one

release concern.

Narrative Review

- Is the thesis clear in the first page?
- Does each section earn its length?
- Are repeated ideas tightened rather than restated?
- Are claims linked to evidence, docs, or clearly named future work?
- Does the paper stay readable for a broad technical or builder audience?

HTML Review

- Does the hosted page read well on desktop and mobile?
- Does the hero lead the page clearly before the supporting surfaces?
- Are diagrams, tables, and screenshots legible without awkward empty runs?
- Do callouts render intentionally instead of leaking raw Markdown markers?
- Are deeper-detail links obvious without overwhelming the page?

PDF Review

- Does the PDF show the white-paper version and updated date clearly?
- Are page breaks reasonable around diagrams, images, tables, and callouts?

- Are any sections visually repetitive or too sparse when printed?
- Is the print styling readable and professional rather than a dark-theme dump?
- Do diagrams and screenshots stay large enough to be useful?

Release Review

- Run `npm run white-paper:review` for the active dev draft.
- Run `npm run white-paper:review:beta` or

`npm run white-paper:review:production` before publishing those lanes.

- Treat the review command as responsible for refreshing the code-review packet

and passing the review gate, not only for rendering a PDF.

- Confirm the review pass also checks preserved-source integrity so active docs

and timing/audio reference work do not drift back to stale machine paths.

- Commit the refreshed review packet before the clean-tree publish step.
- Confirm the generated PDF metadata matches the white-paper version/date.
- Verify the live lane carries both `white-paper.html` and `white-paper.pdf`

after publish.

Related Work Review

- If outside work materially shaped the draft, is it logged?
- If a web refresh was done, are the selected items high-signal and relevant?
- Is the related-work section in the main paper still concise?
- Does the deeper log explain relevance instead of just listing links?

Review Cadence

Small recurring rhythm for when to run the white-paper review spine and what it should catch beyond formatting alone.

Generated from `white-paper/REVIEW_CADENCE.md` during build.

This note turns the reviewer mentality into a small recurring operating rhythm instead of leaving it as a good intention.

Default Cadence

On every meaningful white-paper or preserved-source change

- `run npm run white-paper:review`
- treat the result as a release-surface check, not just a PDF render

Before `/beta` or `/production` publication

- `run npm run white-paper:review:beta` or

`npm run white-paper:review:production`

- verify the lane carries both `white-paper.html` and `white-paper.pdf`

During active fast-moving project periods

- do at least one focused review pass per week on:
 - white-paper clarity
 - HTML/PDF presentation
 - preserved-source integrity
 - stale path drift
 - related-work freshness

On major narrative or methodology shifts

- consider cutting a new white-paper snapshot under `white-paper/releases/`
- refresh the citation ledger and related-work log in the same pass

What The Review Pass Should Catch

- repeated or softened claims that can be tightened
- diagrams, tables, or screenshots that create awkward whitespace
- PDF page breaks that weaken readability
- stale machine-specific source paths in active evidence docs
- preserved-source drift that would make analysis less rerunnable
- related-work references that have gone stale or no longer feel central

Current Review Spine

- `npm run white-paper:review`
 - refresh the dev-lane PDF
 - refresh the code-review packet
 - pass the code-review gate
 - verify the preserved-source integrity check
 - verify the hosted white-paper presentation check

White Paper Working Area

How the project area is organized and when new narrative snapshots should be cut.

Generated from `white-paper/README.md` during build.

This directory holds the durable support files for the project white paper.

Purpose

The white paper is meant to be both:

- a maintained promotional explanation of what Platinum, Aurora, Galaxy

Guardians, ingestion, harnessing, conformance, and release discipline are trying to achieve

- a durable internal reminder of what the project is doing, why it is doing it,

and how the method evolves over time

Structure

- `../WHITE_PAPER.md`
 - current living draft
- `CITATION_LEDGER.md`
 - maintained list of outside ideas, source families, and project learnings
- `ILLUSTRATION_PLAN.md`
 - selected visuals, candidate deeper assets, and open illustration decisions
- `RELATED_WORK.md`
 - dated log of external/public work, why it matters here, and where it fits
- `REVIEWER_CHECKLIST.md`
 - release-minded checklist for HTML, PDF, narrative, and citation review
- `REVIEW_CADENCE.md`
 - the small recurring rhythm for when and how the white-paper review pass

should be run

- `releases/<date>-v<version>/WHITE_PAPER.md`
 - exact snapshot of the white paper for a given release
- `releases/<date>-v<version>/WHITE_PAPER.pdf`
 - printable/exportable snapshot for that white-paper release when available
- `releases/<date>-v<version>/WHITE_PAPER_PDF_METADATA.json`
 - version/date/build metadata for the release PDF snapshot when available
- `releases/<date>-v<version>/RELEASE_NOTES.md`
 - short explanation of what changed in that white paper release

Maintenance Rules

- Update `../WHITE_PAPER.md` directly while shaping the next narrative draft.
- Cut a new snapshot when the white paper changes in a strategically meaningful

way.

- Prefer one snapshot per meaningful narrative release rather than one snapshot

per tiny wording edit.

- When cutting a snapshot, preserve the Markdown, the generated PDF, and the

PDF metadata together whenever the release PDF exists.

- Keep the citation ledger current whenever outside work, source material, or

methodological influences materially shape the paper.

- Keep the illustration plan current whenever a new screenshot, chart, diagram,

or hosted-detail surface becomes important to the white-paper story.

- Keep the related-work log current when a focused web refresh produces a

source worth preserving for future readers.

- Treat HTML and PDF presentation review as part of release quality, not only as

a late formatting pass.

Maintained Release Surfaces

- `dist/<lane>/white-paper.html`
 - hosted readable white-paper surface for a release lane
- `dist/<lane>/white-paper.pdf`
 - printable/exportable white-paper surface for a release lane
- `dist/<lane>/white-paper-pdf.json`
 - version, date-updated, lane, and build metadata for the generated PDF

Suggested Review Commands

- `npm run white-paper:review`
 - refresh the dev-lane PDF, reviewer packet, review gate,

preserved-source-integrity check, and presentation checks together

- `npm run white-paper:review:beta`
 - run the same white-paper review spine for the beta lane
- `npm run white-paper:review:production`
 - run the same white-paper review spine for the production lane

Suggested White Paper Release Triggers

- major public release family change
- major conformance-program reframing
- significant Galaxy Guardians maturity change
- meaningful new evidence or economics story
- recovery of an important prior assessment or external citation

Current Snapshot

The first seeded snapshot is:

- `releases/2026-05-16-v0.1.0/`

This establishes the initial narrative baseline, the release structure, and the citation ledger.

White Paper v0.1.0 Release Notes

Release note for the first seeded white-paper snapshot.

Release: 2026-05-16-v0.1.0 Date: 2026-05-16 Type: foundational narrative release

Summary

This is the first seeded release of the project white paper.

It establishes:

- a single maintained narrative for Platinum, Aurora, Galaxy Guardians, ingestion, harnessing, conformance, release discipline, and AI-assisted engineering
- a dedicated citation ledger
- a release structure for preserving historical white paper versions over time

What This Release Adds

- a preamble explaining the broader project claim
- a section-by-section first draft of the white paper
- diagrams and visual anchors for the operating model
- an initial timeline of major project phases
- a clear note that the earlier Karpathy-style assessment still needs to be

recovered and linked directly

Recommended Next Pass

- tighten the audience voice depending on whether the next target is public

promotion, collaborator onboarding, or AI-method storytelling

- add selected screenshots or dashboard graphics
- recover the missing prior citation material
- revise the paper alongside the next strategic release or conformance shift